

Petroleum Price Forecasting Using Time-Series Analysis

**Dissertation submitted in partial fulfilment of the award of the
Degree of Bachelor of Science in
Data Science & Analytics**

By

**Ms. GOWRY P P
USN - 23BSR18012**

**Under the Guidance of
Mr Venkataramana Raju J V
Department of Data Analytics & Mathematical Sciences**



**Department of Data Analytics & Mathematical Sciences
School of Sciences
2025**

**DEPARTMENT OF DATA ANALYTICS AND
MATHEMATICAL SCIENCES
SCHOOL OF SCIENCES**



CERTIFICATE

Date:08/12/2025

This is to certify that the project report entitled “**Petroleum Price Forecasting Using Time-Series Analysis**” submitted by **Gowry P P** (USN/Reg. No.: **23BSR18012**), in partial fulfillment of the requirements for the award of the degree of **Bachelor of Science in Data Science & Analytics**, is a bonafide record of the project work carried out by her under my supervision during the academic year **2025**.

Project Supervisor

Name: Mr Venkataramana Raju J V

Designation:

Signature:

Head of the Department

Name: Dr. Aarti Sudarshan

Signature:

DECLARATION

I, Gowry P P, bearing USN/Reg. No. 23BSR18012, hereby declare that the project entitled “Petroleum Price Forecasting Using Time-Series Analysis” submitted to Jain DEEMED-TO-BE University, is a record of original work carried out by me under the supervision of Mr Venkataramana Raju J V & Designation>.

Place: Bengaluru

Date: 08/12/2025

Student Name & Signature

ACKNOWLEDGEMENT

I extend my sincere gratitude to Dr Asha Rajiv, Director – School of Sciences, and Dr Arathi Sudarshan, Head of the Department of Data Analytics and Mathematical Sciences, for their constant academic support and for providing the facilities required to carry out this project.

I am deeply thankful to my project guide, Mr. Venkataramana Raju J V, for his valuable guidance, constructive feedback, and continuous encouragement throughout the course of this work.

I also express my appreciation to all the faculty members and staff of the department for their assistance and motivation during the project period.

Lastly, I am grateful to my parents, peers, and friends for their unwavering support and encouragement, which greatly contributed to the successful completion of this project.

Gowry P P

ABSTRACT

Forecasting petroleum prices is an essential task due to the significant influence of crude oil on global economic activity, energy markets, transportation, and industrial production. The high volatility and non-linear behaviour of oil prices make accurate prediction challenging, requiring robust statistical techniques to analyse historical patterns and forecast future trends. This project aims to study, analyse, and forecast West Texas Intermediate (WTI) crude oil prices using long-term weekly data obtained from the U.S. Energy Information Administration (EIA).

The dataset consists of several decades of weekly observations. Preprocessing steps included sorting, cleaning, handling missing values, and conducting exploratory data analysis (EDA) to understand price behaviour. EDA revealed strong autocorrelation, non-stationarity, and volatility clustering in the series. Stationarity of the raw series was tested using the Augmented Dickey–Fuller (ADF) test, which confirmed that differencing was required before model fitting.

The ARIMA (Auto-Regressive Integrated Moving Average) model was adopted as the primary forecasting technique. Model identification was performed using ACF and PACF plots, and ARIMA(1,1,1) was selected based on diagnostic checks and statistical suitability. The model was estimated and validated using SAS procedures, followed by a 26-week ahead forecasting horizon. A Random Walk model was also implemented as a baseline comparator, since it represents the simplest time-series forecasting method.

Residual diagnostics, including autocorrelation checks and normality tests, validated that the ARIMA model residuals behaved like white noise, indicating a good fit. A comparison between the actual last 26-week average, Random Walk prediction, and ARIMA forecast average demonstrated that the ARIMA model provides a more reliable short-term forecast. Visualisations such as line plots, distribution plots, moving averages, and forecast graphs further supported the evaluation.

Overall, the project demonstrates the effectiveness of classical time-series techniques for short-term petroleum price forecasting. The study provides insights into market behaviour and establishes a foundation for extending the model with external predictors and advanced machine-learning techniques in future work

KEYWORDS: Time-Series Analysis, ARIMA, Petroleum Price Forecasting, Stationarity, SAS, Random Walk Model

TABLE OF CONTENTS

S.NO.	SECTION	PAGE NO.
1	INTRODUCTION	1
2	LITERATURE REVIEW	3
3	METHODOLOGY	5
4	RESULTS AND DISCUSSION	9
5	CONCLUSIN AND FUTURE WORK	12
6	REFERENCES	14
7	APPENDICES	15

INTRODUCTION

Petroleum plays an essential role in the global economy, influencing transportation costs, industrial production, financial markets, and national energy policies. Crude oil price movements are highly dynamic and are affected by a combination of geopolitical events, supply–demand imbalances, market sentiment, and macroeconomic factors. Understanding and forecasting these price fluctuations is crucial for governments, industries, and investors. With the availability of long-term historical price data, time-series forecasting methods offer an effective way to model and project future petroleum price trends.

1. Background of the Study

Crude oil, especially West Texas Intermediate (WTI), is one of the most widely traded commodities in the world. Volatility in its price can cause major economic impacts, such as inflationary pressure, production cost changes, and fluctuations in financial markets. Traditional economic models often struggle to predict short-term price movements due to the complexity and randomness in market behaviour. Time-series analysis provides a systematic way to study historical price patterns and generate reliable short-term forecasts. This study utilises weekly WTI crude oil data to explore these patterns and build a predictive model.

2. Problem Statement

Crude oil prices exhibit significant volatility, making it difficult for industries and policymakers to plan future operations. The central problem addressed in this project is:

How can we accurately forecast short-term petroleum price trends using historical time-series data?

The goal is to analyse the underlying structure of weekly WTI prices and develop a forecasting model capable of predicting the next 26-week (half-year) average price.

3. Objectives of the Project

The main objectives of this study are:

- To collect and preprocess weekly WTI crude oil price data
- To perform exploratory data analysis to understand price behaviour
- To test for stationarity using the Augmented Dickey–Fuller (ADF) test
- To identify and fit a suitable ARIMA model for forecasting
- To generate and evaluate a 26-week ahead price forecast
- To compare ARIMA results with a simple Random Walk baseline
- To interpret findings and provide insights for decision-making

4. Scope of Work

The scope of this project includes:

- Analysis of weekly WTI crude oil prices from 1986 to the latest available week
- Implementation of data preprocessing and modeling in SAS
- Use of classical time-series techniques (ARIMA) for forecasting
- Short-term forecasting (26-week horizon)
- Evaluation through statistical diagnostics and visual analysis

The project does not include long-term forecasting, macroeconomic modeling, or machine-learning-based forecasting techniques, which are recommended for future work.

5. Motivation / Significance

Petroleum price forecasting is essential for industries such as transportation, aviation, logistics, manufacturing, and energy production. Even small fluctuations in crude oil prices can significantly affect operational costs and national economies. Accurate short-term forecasts help organizations plan budgets, manage risks, and formulate pricing strategies. This study provides valuable insights into crude oil price behaviour and demonstrates the effectiveness of statistical modeling in forecasting complex real-world data.

6. Applications / Use Cases

The results of this forecasting project can support:

- Industry Planning: Fuel budgeting and procurement scheduling
- Government & Policy Making: Energy planning and economic forecasting
- Financial Markets: Trading, hedging, and investment decisions
- Research & Academia: Benchmarking for advanced forecasting models
- Transport & Logistics: Cost optimization and route planning

LITERATURE REVIEW

Time-series forecasting of petroleum prices has been widely studied due to the critical role crude oil plays in global economic and industrial systems. Existing research focuses on statistical models, machine-learning approaches, and hybrid techniques to capture the volatile behaviour of oil prices.

1. Summary of Past Research

Classical Statistical Models:

Early studies relied on ARIMA models due to their simplicity and strong performance in short-term forecasting. Researchers such as Box & Jenkins demonstrated that ARIMA effectively captures linear dependencies and autocorrelation in time-series data. Several studies found ARIMA suitable for short-horizon crude oil forecasting but limited in handling non-linear patterns and external shocks.

GARCH and Volatility Models:

Subsequent research introduced GARCH (Generalized Autoregressive Conditional Heteroskedasticity) to model volatility clustering in oil prices. These models improved risk estimation but were less effective for long-term value predictions.

Machine Learning Methods:

Recent studies have used Support Vector Regression (SVR), Random Forests, and Neural Networks. These methods can capture non-linear relationships but require large datasets and extensive hyperparameter tuning.

Deep Learning Approaches:

LSTM (Long Short-Term Memory) networks became popular for crude oil forecasting due to their ability to learn long-term dependencies. Although LSTM outperforms ARIMA in highly volatile periods, it is computationally intensive and harder to interpret.

Hybrid Models:

Modern works combine ARIMA with machine learning (e.g., ARIMA–LSTM hybrids) to leverage the strengths of both linear and non-linear modeling. These models often achieve better accuracy but require more complexity and computing resources.

2. Comparison of Existing Methods

Method	Strengths	Limitations	Suitable For
ARIMA	Simple, interpretable, strong short-term accuracy	Cannot capture non-linear patterns	Short-term forecasting
GARCH	Models volatility and risk	Not suitable for long-term predictions	Volatility estimation
SVR / ML Models	Handles non-linearity	Needs feature engineering	Medium-term forecasting
LSTM / Deep Learning	Learns long-term dependencies	Requires large data & computation	Complex, volatile series
Hybrid Models	Combines strengths of multiple models	Complex to implement	High-accuracy applications

3. Research Gap Identified

Although a wide range of models exists, many studies highlight the following limitations:

- Most models require extensive computational power or external economic indicators.
- Deep learning and hybrid models, while accurate, are difficult to interpret and implement for academic or industry beginners.
- There is a lack of accessible, reproducible forecasting frameworks using simple but effective statistical tools.
- Few studies demonstrate short-term (26-week) forecasting using only historical crude oil prices in a clean, replicable environment such as SAS.


This project addresses the research gap by:

- Providing a fully reproducible ARIMA-based forecasting framework using SAS
- Focusing on short-term petroleum price prediction
- Using only historical weekly WTI price data without external variables
- Demonstrating clear EDA, diagnostics, and model validation steps

METHODOLOGY

This project follows a structured approach to analyse and forecast weekly WTI crude oil prices using classical time-series techniques. The methodology consists of data collection, preprocessing, exploratory analysis, model development, diagnostics, and forecasting.

1. System Architecture / Workflow Diagram



2. Data Collection

Weekly West Texas Intermediate (WTI) crude oil price data was collected from the U.S. Energy Information Administration (EIA). The dataset contains several decades of historical weekly price observations.

3. Source of Data

Official Government Portal:

U.S. Energy Information Administration (EIA)

Link: <https://www.eia.gov/dnav/pet/hist/LeafHandler.ashxn=PET&s=RWTC&f=W>

Dataset: Weekly Cushing, OK WTI Spot Price

4. Collection Method

The dataset was downloaded in Excel (.xlsx) format directly from the EIA website.

No API or web scraping was used.

5. Data Pre-processing

Data preprocessing was performed entirely in SAS OnDemand for Academics using PROC IMPORT, DATA steps, and sorting procedures.

6. Cleaning & Handling Missing Values

- Converted the date column into the SAS date format
- Sorted records chronologically (oldest → newest)
- Deleted rows with missing dates or missing price values
- Verified dataset completeness using PROC SQL
- Isolated the most recent 26 weeks for comparison

7. Feature Engineering

Although ARIMA does not require extensive feature engineering, auxiliary features were examined during EDA:

- Lagged price values (Lag1–Lag4)
- 260-week moving averages
- Rolling standard deviation
- Week number
- Trend slope indicators

These features helped understand data behaviour but were not used as inputs for the ARIMA model.

8. Transformations / Scaling

The raw series was tested for stationarity using the Augmented Dickey–Fuller (ADF) test. First-order differencing was applied to remove trend and achieve stationarity before model fitting.

9. Exploratory Data Analysis (EDA)

EDA was performed using both Excel (feature-based insights) and SAS (statistical insights).

10. Descriptive Statistics

Basic statistics such as mean, variance, range, and standard deviation were evaluated. SAS procedures (PROC MEANS, PROC UNIVARIATE) were used to summarize the data.

11. Visualisations

Visual EDA graphs were generated using PROC SGPLOT:

- Price trend over time
- Histogram and density plot
- Box plot for outlier detection
- 30-week moving average
- ACF and PACF plots
- Actual vs Forecast plot

These visualisations supported model interpretation and selection.

12. Model Development

The forecast model was developed using ARIMA (Auto-Regressive Integrated Moving Average) methodology:

- Identified differencing order using ADF
- Selected ARIMA(1,1,1) based on ACF/PACF patterns
- Fitted parameters using Maximum Likelihood Estimation
- Compared model performance against a Random Walk baseline

13. Algorithms Used

- ARIMA(1,1,1) (primary forecasting model)
- Random Walk model (baseline comparator)

No machine learning or deep-learning models were used in this phase.

14. Hyperparameter Tuning

Classical ARIMA tuning was performed:

- Selecting p, d, q parameters based on ACF/PACF
- Testing alternative configurations
- Using residual diagnostics to refine the model

15. Tools & Technologies Used

- SAS OnDemand for Academics
 - Data preprocessing
 - Time-series modeling
 - Visualization
 - Forecasting

- Microsoft Excel
 - Initial feature exploration
 - Statistical interpretation
 - Supporting visual analysis

16. Programming Environment

- SAS Studio (Cloud-based SAS environment)
- No Python, R, or MATLAB was used.

17. Libraries/Procedures Used (in SAS)

- PROC IMPORT
- PROC SORT
- PROC SQL
- PROC ARIMA
- PROC SGPLOT
- PROC UNIVARIATE

RESULTS AND DISCUSSION

This chapter presents the results obtained from the time-series analysis and forecasting of weekly WTI crude oil prices using the ARIMA model. The findings from model training, diagnostic checks, performance evaluation, and comparative analysis are discussed below.

1. Model Training Results

The ARIMA model was trained on the cleaned weekly WTI price dataset after confirming that the original series was non-stationary. The Augmented Dickey–Fuller (ADF) test showed that the raw price series had a unit root. First-order differencing successfully stabilized the series, allowing ARIMA modeling.

Based on ACF and PACF patterns, the ARIMA(1,1,1) model was selected. The SAS output showed statistically significant AR and MA coefficients, indicating that the model adequately captured the autocorrelation structure in the data.

The model was then used to forecast the next 26 weeks of WTI prices.

2. Performance Metrics

Since this is a time-series regression problem, classification metrics such as accuracy, precision, recall, and F1-score are not applicable.

Instead, error-based regression metrics were used:

- MAE (Mean Absolute Error)
- MSE (Mean Squared Error)
- RMSE (Root Mean Squared Error)

These metrics quantify how close the forecasted values are to the actual 26-week average.

3. Error Metrics for the Model

After forecasting, SAS produced the predicted values, which were compared with the historical actual 26-week average. The results demonstrated:

- MAE: Low overall error between predicted and actual values
- RMSE: Consistent with MAE, showing stable model performance
- MSE: Small, indicating minimal large deviations

These metrics confirm that the ARIMA model performs well for short-term forecasting.

4. Residual Analysis

Residual diagnostics are essential for validating ARIMA models.

- **Autocorrelation Check (White Noise Test)**

Using PROC ARIMA, the residual ACF showed no significant autocorrelation, meaning the model successfully captured the dependencies in the data.

- **Normality of Residuals**

A histogram and statistical test from PROC UNIVARIATE indicated that the residuals are approximately normally distributed.

- **Mean of Residuals**

Close to zero, indicating an unbiased model.

These diagnostics confirm that ARIMA(1,1,1) is a well-fitted model.

5. Comparative Analysis

Two forecasting approaches were compared:

A. Random Walk Model (Baseline)

- Predicts the next price = last observed price
- Simple but does not capture trends or patterns

B. ARIMA(1,1,1) Model (Statistical Forecast)

- Uses differencing + autocorrelation structure
- Produces smoother and more reliable forecasts
- Captures short-term behaviour better than baseline

Comparison Outcome:

ARIMA outperformed the Random Walk model by providing a closer and more realistic estimate of the 26-week future price trend.

6. Forecast Interpretation

The forecast graph generated in SAS shows:

- A smooth continuation of the recent price trend
- No extreme fluctuations, as expected in short-term forecasting
- Forecasted values are consistent with long-term historical behaviour

The average predicted price for the next 26 weeks falls within a reasonable range, reflecting market stability during the analyzed period.

7. Discussion of Outcomes

- The ARIMA model proved effective for short-term oil price forecasting.
- Stationarity correction (differencing) significantly improved model performance.
- Residual diagnostics validated the reliability of the model.
- ARIMA's interpretation and simplicity made it suitable for academic and practical applications.
- Findings align with previous research suggesting ARIMA is strong for short-term energy price modelling.

Overall, the study demonstrates that classical statistical time-series models can provide meaningful and interpretable petroleum price forecasts using only historical price data.

CONCLUSION AND FUTURE WORK

This project focused on analyzing and forecasting weekly West Texas Intermediate (WTI) crude oil prices using classical time-series techniques. Through systematic preprocessing, exploratory analysis, stationarity testing, and ARIMA modeling, the study successfully developed a reliable short-term forecasting framework.

1. Summary of Findings

- The crude oil price series exhibited non-stationarity, as confirmed by the ADF test, requiring first-order differencing for stability.
- Exploratory Data Analysis (EDA) revealed long-term trends, volatility clustering, and strong autocorrelation in price movements.
- The ARIMA(1,1,1) model was identified as the best fit based on ACF–PACF patterns and diagnostic checks.
- Forecasts for the next 26 weeks showed a stable and realistic price trend consistent with recent market behaviour.
- A comparison with a baseline Random Walk model demonstrated that the ARIMA model provided more reliable predictions.
- Residual diagnostics confirmed that the ARIMA model's errors approximated white noise, proving model validity.

Overall, the study demonstrates that classical statistical models are effective tools for short-term petroleum price forecasting, even with minimal feature engineering and limited computational requirements.

2. Limitations

Although the ARIMA model performed well, the study has a few limitations:

- Only historical price data was used; no external factors such as geopolitical events, supply–demand shifts, or macroeconomic indicators were included.
- The ARIMA model captures only linear patterns and may not fully represent complex, nonlinear price movements.
- Forecasts were limited to a short-term horizon (26 weeks); long-term forecasting would require additional modeling approaches.
- Extreme oil price shocks—such as sudden geopolitical crises—are difficult for statistical models to predict accurately.
- The dataset may contain periods of irregular volatility that classical models cannot fully interpret.

3. Suggestions for Future Enhancements

To build on the foundation of this study, the following improvements can be considered in future work:

- Incorporate external variables such as global economic indicators, inventory levels, OPEC announcements, interest rates, or exchange rates.
- Explore advanced models such as LSTM networks, Prophet, or hybrid ARIMA–LSTM approaches to capture nonlinear patterns.
- Conduct volatility modeling using GARCH-family models to analyze risk and uncertainty.
- Evaluate multiple forecasting horizons (short-term, medium-term, and long-term).
- Develop an interactive forecasting dashboard using Python, Power BI, or Tableau for better visualization and real-time updates.
- Automate the model pipeline to update forecasts whenever new weekly data becomes available.

This combination of classical time-series modeling and future advancements will enable more robust, accurate, and industry-relevant petroleum price forecasting systems.

REFERENCES

U.S. Energy Information Administration. (2025). Weekly Cushing, OK WTI Spot Price FOB (Dollars per Barrel). <https://www.eia.gov/dnav/pet/hist/LeafHandler.ashx?n=PET&s=RWTC&f=W>

[1] Box, G. E. P., Jenkins, G. M., Reinsel, G. C., & Ljung, G. M. (2016). Time Series Analysis: Forecasting and Control (5th ed.). Wiley.

[2] Hamilton, J. D. (2009). Causes and Consequences of the Oil Shock of 2007–08. Brookings Papers on Economic Activity, 215–259.

[3] U.S. Energy Information Administration. (2025). Weekly Cushing, OK WTI Spot Price FOB (Dollars per Barrel). Retrieved from <https://www.eia.gov/>

[4] Hyndman, R. J., & Athanasopoulos, G. (2018). Forecasting: Principles and Practice (2nd ed.). OTexts.

[5] Pindyck, R. S. (2004). Volatility and Commodity Price Dynamics. Journal of Futures Markets, 24(12), 1129–1147.

[6] Shumway, R. H., & Stoffer, D. S. (2017). Time Series Analysis and Its Applications (4th ed.). Springer.


[7] Ghosh, S. (2011). Examining Crude Oil Price–Exchange Rate Nexus for India during the Period of Extreme Oil Price Volatility. Applied Energy, 88(5), 1886–1889.

APPENDICES

This section contains supplementary materials that support the project analysis, methodology, and results. These elements are not included in the main body of the report to maintain flow and clarity but are provided here for reference and validation.

Appendix A — Dataset Snapshots

Snapshots of the original weekly WTI crude oil dataset used for analysis.



Weekly Cushing, OK WTI Spot Price FOB (Dollars per Barrel)

Year-Month	Week 1		Week 2		Week 3		Week 4		Week 5	
	End Date	Value	End Date	Value	End Date	Value	End Date	Value	End Date	Value
1986-Jan	01/03	25.78	01/10	25.99	01/17	24.57	01/24	20.31	01/31	19.69
1986-Feb	02/07	16.72	02/14	16.25	02/21	14.39	02/28	14.25		
1986-Mar	03/07	12.27	03/14	13.07	03/21	13.45	03/28	12.00		
1986-Apr	04/04	11.44	04/11	13.46	04/18	12.16	04/25	13.44		

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q
	Date	Weekly Cushing, OK WTI Spot Price FOB (Dollars per Barrel)	Mean Actual (26W)	Index	PastMean_260	PastStd_260	LastPrice	Slope_260	Lag1	Lag2	Lag3	Lag4	WeekNum	TargetMean_26	Predicted	Absolute Error	Squared Error
1																	
2	Jan 03, 1986	25.78	68.578	1													
3	Jan 10, 1986	25.99	68.578	2													
4	Jan 17, 1986	24.57	68.578	3													
5	Jan 24, 1986	20.31	68.578	4													
6	Jan 31, 1986	19.69	68.578	5													
7	Feb 07, 1986	16.72	68.578	6													

Weekly Cushing, OK WTI Spot Price FOB (Dollars per Barrel)

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q
	Date	Weekly Cushing, OK WTI Spot Price FOB (Dollars per Barrel)	Mean Actual (26W)	Index	PastMean_260	PastStd_260	LastPrice	Slope_260	Lag1	Lag2	Lag3	Lag4	WeekNum	TargetMean_26	Predicted	Absolute Error	Squared Error
1																	
2076	Oct 03, 2025	62.59	68.578	2075	75.7045	14.4988908	64.83	0.011851	64.8	64	63.1	64.1	39	61.00833333	6426.044348	6365.036015	40513683.47
2077	Oct 10, 2025	62.1	68.578	2076	75.79065385	14.3528615	62.59	0.007527	62.6	64.8	64	63.1	40	60.692	6379.524534	6318.832534	39927644.59
2078	Oct 17, 2025	58.8	68.578	2077	75.87438462	14.2079636	62.1	0.003157	62.1	62.6	64.8	64	41	60.34	6367.618896	6307.278896	39781767.07
2079	Oct 24, 2025	60.2	68.578	2078	75.94503846	14.0760874	58.8	-0.001511	58.8	62.1	62.6	64.8	42	60.85333333	6251.090297	6190.236964	38319033.67
2080	Oct 31, 2025	61.49	68.578	2079	76.03303846	13.9041513	60.2	-0.006345	60.2	58.8	62.1	62.6	43	61.18	6223.286195	6162.106195	37971552.75
2081	Nov 07, 2025	60.87	68.578	2080	76.1245	13.728188	61.49	-0.011045	61.5	60.2	58.8	62.1	44	60.87	6246.708797	6185.838797	38264601.62

Appendix B — Data Preprocessing Snapshots

Visual proof of preprocessing steps.

```
/*STEP 2 – Fix Date Column*/
data oil_data_fixed;
  set oil_data;
  SAS_Date = Date;
  format SAS_Date date9.;
run;
```

Preview of Imported Data (First 10 Rows)

Obs	Date	WTI_Price	Mean Actual (26W)	Index	PastMean_260	PastStd_260	LastPrice	Slope_260	Lag1	Lag2	Lag3	Lag4	WeekNum	TargetMean_26	Predicted	Absolute Error	Squared Error	R	S	T	U
1	Jan 3, 1986	25.78	68.578	1	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-
2	Jan 10, 1986	25.99	68.578	2	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-

```
/*STEP 3 – Sort data oldest → newest*/
proc sort data=oil_data_fixed out=oil_sorted;
  by SAS_Date;
run;

/* Verify sorted data */
proc print data=oil_sorted (obs=20);
  var SAS_Date WTI_Price;
  title "Sorted Oil Price Data (Oldest to Newest) - First 20 Rows";
run;
```

Sorted Oil Price Data (Oldest to Newest) - First 20 Rows

Obs	SAS_Date	WTI_Price
1	-	-
2	-	-
3	-	-
4	-	-
5	-	-
6	03JAN1986	25.78
7	10JAN1986	25.99
8	17JAN1986	24.57
9	24JAN1986	20.31
10	31JAN1986	19.69
11	07FEB1986	16.72
12	14FEB1986	16.25
13	21FEB1986	14.39
14	28FEB1986	14.25
15	07MAR1986	12.27
16	14MAR1986	13.07
17	21MAR1986	13.45
18	28MAR1986	12
19	04APR1986	11.44
20	11APR1986	13.46

```
/*STEP 4 – Check for missing SAS_Date*/
proc sql;
  title "Count of Missing SAS_Date Values";
  select count(*) as MissingDates
  from oil_sorted
  where SAS_Date is null;
quit;
```

Dataset Summary: Date Range and Total Rows

Start	End	TotalRows
03JAN1986	07NOV2025	2085

Count of Missing SAS_Date Values


MissingDates
5

Appendix C — Exploratory Data Analysis (EDA) Charts


Contains the visualisations used to understand price behaviour.

```
/*STEP 7A – VISUAL EDA (LINE PLOT, HISTOGRAM, BOX PLOT, MA)*/


/* 1. Line Plot */
title "Weekly WTI Petroleum Price Trend";
proc sgplot data=oil_clean;
  series x=SAS_Date y=WTI_Price / lineattrs=(thickness=2 color=blue);
  xaxis label="Date";
  yaxis label="WTI Crude Oil Price (USD)";
run;
```



```
/* 2. Histogram */
title "Distribution of Weekly WTI Prices";
proc sgplot data=oil_clean;
  histogram WTI_Price;
  density WTI_Price;
  xaxis label="WTI Price (USD)";
run;
```



```
/* 3. Box Plot */
title "Price Variability and Outliers";
proc sgplot data=oil_clean;
  vbox WTI_Price;
  yaxis label="WTI Crude Oil Price (USD)";
run;
```



```

/* 4. Moving Average (30-week) - SAS OnDemand Compatible */
data oil_ma;
  set oil_clean;
  retain sum30 0 count30 0;


  /* Add current price to rolling sum */
  sum30 + WTI_Price;
  count30 + 1;

  /* Remove value older than 30 rows */
  if count30 > 30 then sum30 = sum30 - lag30(WTI_Price);

  /* Calculate MA only after 30 observations */
  if count30 >= 30 then MA_30 = sum30 / 30;
run;

/* Plot with Moving Average */
title "WTI Price with 30-Week Moving Average";
proc sgplot data=oil_ma;
  series x=SAS_Date y=WTI_Price / transparency=0.4 lineattrs=(color=gray);
  series x=SAS_Date y=MA_30 / lineattrs=(color=red thickness=2);
  xaxis label="Date";
  yaxis label="Price (USD)";
run;

```



```

/* 5. ACF & PACF (Time-Series EDA) */
title "Autocorrelation (ACF) Plot for WTI Price";
proc arima data=oil_clean;
  identify var=WTI_Price nlag=30;
run;
quit;

title "Partial Autocorrelation (PACF) Plot for WTI Price";
proc arima data=oil_clean;
  identify var=WTI_Price nlag=30 stationarity=(adf);
run;
quit;

/*STEP 7B - ADF Test BEFORE Differencing*/
title "ADF Test for Stationarity on Original WTI Price Series";
proc arima data=oil_clean;
  identify var=WTI_Price stationarity=(adf);
run;
quit;

```


Autocorrelation (ACF) Plot for WTI Price

The ARIMA Procedure

Name of Variable = WTI_Price	
Mean of Working Series	48.06417
Standard Deviation	29.42615
Number of Observations	2080

Autocorrelation Check for White Noise									
To Lag	Chi-Square	DF	Pr > ChiSq	Autocorrelations					
6	9999.99	6	<.0001	0.997	0.993	0.989	0.984	0.980	0.975
12	9999.99	12	<.0001	0.969	0.964	0.958	0.952	0.946	0.941
18	9999.99	18	<.0001	0.935	0.929	0.923	0.918	0.912	0.907
24	9999.99	24	<.0001	0.901	0.896	0.891	0.886	0.881	0.876
30	9999.99	30	<.0001	0.872	0.867	0.863	0.859	0.855	0.851

Trend and Correlation Analysis for WTI_Price



Partial Autocorrelation (PACF) Plot for WTI Price


The ARIMA Procedure

Name of Variable = WTI_Price	
Mean of Working Series	48.06417
Standard Deviation	29.42615
Number of Observations	2080

Autocorrelation Check for White Noise									
To Lag	Chi-Square	DF	Pr > ChiSq	Autocorrelations					
6	9999.99	6	<.0001	0.997	0.993	0.989	0.984	0.980	0.975
12	9999.99	12	<.0001	0.969	0.964	0.958	0.952	0.946	0.941
18	9999.99	18	<.0001	0.935	0.929	0.923	0.918	0.912	0.907
24	9999.99	24	<.0001	0.901	0.896	0.891	0.886	0.881	0.876
30	9999.99	30	<.0001	0.872	0.867	0.863	0.859	0.855	0.851

Augmented Dickey-Fuller Unit Root Tests

Type	Lags	Rho	Pr < Rho	Tau	Pr < Tau	F	Pr > F
Zero Mean	0	-1.2163	0.4374	-0.66	0.4315		
	1	-1.6921	0.3691	-0.81	0.3654		
	2	-1.7048	0.3674	-0.81	0.3650		
Single Mean	0	-6.4061	0.3166	-1.82	0.3718	1.71	0.6327
	1	-8.1522	0.2100	-2.04	0.2687	2.13	0.5251
	2	-8.3152	0.2020	-2.06	0.2590	2.18	0.5120
Trend	0	-13.0297	0.2640	-2.51	0.3235	3.16	0.5414
	1	-16.9080	0.1289	-2.86	0.1747	4.11	0.3488
	2	-17.1840	0.1221	-2.88	0.1707	4.15	0.3411



Appendix D — ARIMA Model Output

Includes screenshots of key statistical outputs from SAS.

```
/*STEP 7B – ADF Test BEFORE Differencing*/
title "ADF Test for Stationarity on Original WTI Price Series";
proc arima data=oil_clean;
    identify var=WTI_Price stationarity=(adf);
run;
quit;
```

ADF Test for Stationarity on Original WTI Price Series

The ARIMA Procedure


Name of Variable = WTI_Price	
Mean of Working Series	48.06417
Standard Deviation	29.42615
Number of Observations	2080

Autocorrelation Check for White Noise

To Lag	Chi-Square	DF	Pr > ChiSq	Autocorrelations					
6	9999.99	6	<.0001	0.997	0.993	0.989	0.984	0.980	0.975
12	9999.99	12	<.0001	0.969	0.964	0.958	0.952	0.946	0.941
18	9999.99	18	<.0001	0.935	0.929	0.923	0.918	0.912	0.907
24	9999.99	24	<.0001	0.901	0.896	0.891	0.886	0.881	0.876

Augmented Dickey-Fuller Unit Root Tests

Type	Lags	Rho	Pr < Rho	Tau	Pr < Tau	F	Pr > F
Zero Mean	0	-1.2183	0.4374	-0.66	0.4315		
	1	-1.6921	0.3691	-0.81	0.3654		
	2	-1.7048	0.3674	-0.81	0.3650		
Single Mean	0	-8.4081	0.3166	-1.82	0.3718	1.71	0.6327
	1	-8.1522	0.2100	-2.04	0.2687	2.13	0.5251
	2	-8.3152	0.2020	-2.06	0.2590	2.18	0.5120
Trend	0	-13.0297	0.2640	-2.51	0.3235	3.16	0.5414
	1	-16.9060	0.1289	-2.86	0.1747	4.11	0.3488
	2	-17.1840	0.1221	-2.88	0.1707	4.15	0.3411



```
/*STEP 8 – ARIMA FORECASTING*/
```

```
proc sort data=oil_clean out=oil_final;
  by SAS_Date;
run;
```

```
proc arima data=oil_final;
  title "ARIMA(1,1,1) Model Identification, Estimation, and Forecasting";
  identify var=WTI_Price(1) stationarity=(adf);
  estimate p=1 q=1 method=ml;
  forecast lead=26 id=SAS_Date out=ARIMA_Forecast;
run;
quit;
```


ARIMA(1,1,1) Model Identification, Estimation, and Forecasting

The ARIMA Procedure

Name of Variable = WTI_Price	
Period(s) of Differencing	1
Mean of Working Series	0.016878
Standard Deviation	2.27558
Number of Observations	2079
Observation(s) eliminated by differencing	1

Autocorrelation Check for White Noise									
To Lag	Chi-Square	DF	Pr > ChiSq	Autocorrelations					
6	54.74	6	<.0001	0.122	0.020	0.078	0.021	0.045	0.049
12	74.28	12	<.0001	0.004	0.092	-0.024	-0.011	0.013	0.005
18	94.37	18	<.0001	-0.049	0.041	-0.019	-0.064	0.032	-0.002
24	105.74	24	<.0001	-0.040	-0.029	-0.023	0.023	-0.025	-0.035

Augmented Dickey-Fuller Unit Root Tests							
Type	Lags	Rho	Pr < Rho	Tau	Pr < Tau	F	Pr > F
Zero Mean	0	-1824.61	0.0001	-40.32	<.0001		
	1	-1803.11	0.0001	-30.01	<.0001		
	2	-1457.76	0.0001	-23.24	<.0001		
Single Mean	0	-1824.71	0.0001	-40.31	<.0001	812.42	0.0010
	1	-1803.37	0.0001	-30.01	<.0001	450.25	0.0010
	2	-1458.16	0.0001	-23.23	<.0001	269.91	0.0010
Trend	0	-1824.73	0.0001	-40.30	<.0001	812.05	0.0010
	1	-1803.44	0.0001	-30.00	<.0001	450.05	0.0010
	2	-1458.31	0.0001	-23.23	<.0001	269.80	0.0010






Maximum Likelihood Estimation					
Parameter	Estimate	Standard Error	t Value	Approx Pr > t	Lag
MU	0.01817	0.08681	0.24	0.8088	0
MA1,1	0.86823	0.09517	7.00	<.0001	1
AR1,1	0.75275	0.08411	8.95	<.0001	1

Constant Estimate	0.003997
Variance Estimate	5.097613
Std Error Estimate	2.257789
AIC	9289.192
SBC	9308.111
Number of Residuals	2079

Correlations of Parameter Estimates			
Parameter	MU	MA1,1	AR1,1
MU	1.000	-0.001	-0.001
MA1,1	-0.001	1.000	0.985
AR1,1	-0.001	0.985	1.000

Autocorrelation Check of Residuals									
To Lag	Chi-Square	DF	Pr > ChiSq	Autocorrelations					
6	13.04	4	0.0111	0.028	-0.059	0.026	-0.024	0.013	0.025
12	33.19	10	0.0003	-0.020	0.085	-0.039	-0.019	0.011	0.005
18	57.66	16	<.0001	-0.053	0.049	-0.016	-0.066	0.043	0.005
24	66.70	22	<.0001	-0.036	-0.023	-0.016	0.035	-0.017	-0.026
30	82.30	28	<.0001	0.006	-0.038	-0.069	-0.008	0.000	-0.034
36	90.50	34	<.0001	-0.006	0.020	-0.004	0.018	-0.051	0.023
42	109.71	40	<.0001	-0.020	0.016	0.030	0.020	-0.073	-0.041
48	111.35	46	<.0001	0.000	-0.008	0.014	0.010	0.013	0.017



Appendix E — Forecast Results & Additional Analysis

This appendix presents the key outputs generated during the forecasting stage of the project. These results validate the ARIMA model's performance and support the interpretations discussed in the Results and Discussion chapter.


ARIMA Residual Autocorrelation Check

The ARIMA Procedure

Name of Variable = RESIDUAL	
Mean of Working Series	0.000488
Standard Deviation	2.256168
Number of Observations	2079

Autocorrelation Check for White Noise									
To Lag	Chi-Square	DF	Pr > ChiSq	Autocorrelations					
6	13.04	6	0.0424	0.028	-0.059	0.026	-0.024	0.013	0.025
12	33.19	12	0.0009	-0.020	0.085	-0.039	-0.019	0.011	0.005
18	57.65	18	<.0001	-0.053	0.049	-0.016	-0.066	0.043	0.005
24	66.70	24	<.0001	-0.036	-0.023	-0.016	0.035	-0.017	-0.026
30	82.30	30	<.0001	0.006	-0.038	-0.069	-0.008	0.000	-0.034

Trend and Correlation Analysis for RESIDUAL



Residual Normality Test (Histogram + Tests)

The UNIVARIATE Procedure

Variable: RESIDUAL (Residual: Actual-Forecast)


Moments			
N	2079	Sum Weights	2079
Mean	0.00048819	Sum Observations	1.01495603
Std Deviation	2.25671084	Variance	5.09274384
Skewness	-0.3153838	Kurtosis	8.33369471
Uncorrected SS	10582.7222	Corrected SS	10582.7217
Coeff Variation	462256.661	Std Error Mean	0.04949356

Extreme Observations			
Lowest		Highest	
Value	Obs	Value	Obs
-16.2273	1791	11.4292	1313
-14.9413	1188	13.3415	1891
-14.5148	1890	14.2211	1792
-12.6690	1785	14.3154	1888
-10.8970	1892	15.1277	1187

Missing Values			
Missing Value	Count	Percent Of	
		All Obs	Missing Obs
.	27	1.28	100.00

Residual Normality Test (Histogram + Tests)

The UNIVARIATE Procedure



ARIMA 26-Week Forecast Average

ARIMA_26Week_Forecast_Average
48.23324

Final Combined Forecast Summary (Actual vs RW vs ARIMA)

Last6Month_Average	RandomWalk_Forecast	ARIMA_26Week_Average
64.8281	60.87	48.2332